

In [1]:

```
import matplotlib.pyplot as plt
import numpy as np
import os
import pandas as pd
import seaborn as sns
from IPython.display import Image
from scipy import stats
from sklearn.linear_model import LinearRegression
from sklearn.preprocessing import FunctionTransformer
from sklearn.preprocessing import OneHotEncoder
from sklearn.preprocessing import MinMaxScaler
from sklearn.preprocessing import OrdinalEncoder
from sklearn.pipeline import Pipeline
from sklearn.compose import ColumnTransformer
from sklearn.tree import DecisionTreeRegressor
from sklearn.model_selection import train_test_split
from sklearn.neighbors import KNeighborsRegressor
from sklearn.model_selection import train_test_split
from sklearn.neighbors import KNeighborsClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestRegressor
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score
from sklearn.metrics import confusion_matrix
from sklearn.tree import plot_tree
from sklearn.model_selection import GridSearchCV
from sklearn.model_selection import cross_val_score
from sklearn.cluster import KMeans
from sklearn.mixture import GaussianMixture
from sklearn.svm import SVC
from sklearn.naive_bayes import GaussianNB
from sklearn.ensemble import GradientBoostingClassifier
from sklearn.ensemble import BaggingClassifier
from sklearn.naive_bayes import BernoulliNB
from sklearn.neighbors import NearestCentroid
from sklearn.neural_network import MLPClassifier
import time

import warnings
warnings.filterwarnings('ignore')
verbose = False
```


In [2]:

```
vehicles = pd.read_csv("data/dftRoadSafetyData_Vehicles_2018.csv")  
vehicles.head()
```

Out[2]:

	Accident_Index	Vehicle_Reference	Vehicle_Type	Towing_and_Articulation	Vehicle_Manoeuvre
0	2018010080971	1	9	0	18
1	2018010080971	2	8	0	18
2	2018010080973	1	9	0	18
3	2018010080974	1	8	0	7
4	2018010080974	2	9	0	18

5 rows × 23 columns




In [3]:

```
accidents = pd.read_csv("data/dftRoadSafetyData_Accidents_2018.csv")  
accidents.head()
```

Out[3]:

	Accident_Index	Location_Easting_OSGR	Location_Northing_OSGR	Longitude	Latitude	Police_Force
0	2018010080971	529150.0	182270.0	-0.139737	51.524587	101
1	2018010080973	542020.0	184290.0	0.046471	51.539651	101
2	2018010080974	531720.0	182910.0	-0.102474	51.529746	101
3	2018010080981	541450.0	183220.0	0.037828	51.530179	101
4	2018010080982	543580.0	176500.0	0.065781	51.469258	101

5 rows × 32 columns



In [4]:

```
casualties = pd.read_csv("data/dftRoadSafetyData_Casualties_2018.csv")
casualties.head()
```

Out[4]:

	Accident_Index	Vehicle_Reference	Casualty_Reference	Casualty_Class	Sex_of_Casualty	Age
0	2018010080971	1	1	2	2	
1	2018010080971	2	2	1	1	
2	2018010080973	1	1	3	1	
3	2018010080974	1	1	1	1	
4	2018010080981	1	1	1	1	

In [5]:

```
df = vehicles.merge(accidents, how='inner', on='Accident_Index')
df = df.merge(casualties, how='inner', on='Accident_Index')
df.head()
```

Out[5]:

	Accident_Index	Vehicle_Reference_x	Vehicle_Type	Towing_and_Articulation	Vehicle_Manoeuvr
0	2018010135259	1	9	0	-
1	2018010135259	2	9	0	-
2	2018010135261	1	11	0	1
3	2018010135264	1	9	0	1
4	2018010135268	1	9	0	1

5 rows × 69 columns

In [6]:

```
population = pd.read_excel("data/regionalgrossdomesticproductgdplocalauthorities.xlsx", sheet_name=6, header=1, nrows=382)
population = population[["LA name", 2018]]
population = population.rename({2018:"Population"}, axis=1)
population.head()
```

Out[6]:

	LA name	Population
0	Hartlepool	93242
1	Middlesbrough	140545
2	Redcar and Cleveland	136718
3	Stockton-on-Tees	197213
4	Darlington	106566

In [7]:

```
gdp = pd.read_excel("data/regionalgrossdomesticproductgdplocalauthorities.xlsx", sheet_name=7, header=1, nrows=382)
gdp = gdp[["LA name", '20183']]
gdp = gdp.rename({'20183':"GDP"}, axis=1)
gdp.head()
```

Out[7]:

	LA name	GDP
0	Hartlepool	18572
1	Middlesbrough	24103
2	Redcar and Cleveland	15793
3	Stockton-on-Tees	29843
4	Darlington	28866

In [8]:

```
gdp_growth = pd.read_excel("data/regionalgrossdomesticproductgdplocalauthorities.xlsx", sheet_name=13, header=1, nrows=382)
gdp_growth = gdp_growth[["LA name", '20183']]
gdp_growth = gdp_growth.rename({'20183': "GDP Growth"}, axis=1)
gdp_growth.head()
```

Out[8]:

	LA name	GDP Growth
0	Hartlepool	-2.6
1	Middlesbrough	3.5
2	Redcar and Cleveland	2.6
3	Stockton-on-Tees	-4.8
4	Darlington	-6.4

In [9]:

```
# Merge all financial Data
financial = pd.merge(population, gdp, on='LA name').merge(gdp_growth, on='LA name')
financial
```

Out[9]:

	LA name	Population	GDP	GDP Growth
0	Hartlepool	93242	18572	-2.6
1	Middlesbrough	140545	24103	3.5
2	Redcar and Cleveland	136718	15793	2.6
3	Stockton-on-Tees	197213	29843	-4.8
4	Darlington	106566	28866	-6.4
...
377	Lisburn and Castlereagh	144381	25918	3.6
378	Mid and East Antrim	138773	29885	-10.5
379	Mid Ulster	147392	24661	2.9
380	Newry, Mourne and Down	180012	18408	-3.0
381	Ards and North Down	160864	15034	0.1

382 rows × 4 columns

In [10]:

```
vlookup = pd.read_excel("data/variable lookup.xls", sheet_name=5, header=0)
vlookup = vlookup.rename({"label" : "LA name"},axis=1)
vlookup.head()
```

Out[10]:

	code	LA name
0	1	Westminster
1	2	Camden
2	3	Islington
3	4	Hackney
4	5	Tower Hamlets


In [11]:

```
# Add Names for LA
df = pd.merge(df, vlookup, left_on='Local_Authority_(District)', right_on='code')
df.head()
```

Out[11]:

	Accident_Index	Vehicle_Reference_x	Vehicle_Type	Towing_and_Articulation	Vehicle_Manoeuvr
0	2018010135259	1	9	0	-
1	2018010135259	2	9	0	-
2	2018010135309	1	9	0	1
3	2018010135309	2	11	0	1
4	2018010135320	1	1	-1	-

5 rows × 71 columns




In [12]:

```
# merge with financials
df = pd.merge(df, financial, on='LA name')
df.head()
```

Out[12]:

	Accident_Index	Vehicle_Reference_x	Vehicle_Type	Towing_and_Articulation	Vehicle_Manoeuvr
0	2018010135259	1	9	0	-
1	2018010135259	2	9	0	-
2	2018010135309	1	9	0	1
3	2018010135309	2	11	0	1
4	2018010135320	1	1	-1	-

5 rows × 74 columns



In [13]:

```
# Land Usage
land = pd.read_excel("data/Land_Use_England_2017.xlsx", sheet_name=3, header=8, nrows = 326)
land = land.iloc[:,[1,6,11]]
land = land.rename({"Local authority" : "LA name", "Unnamed: 6":"Residential", "Unnamed: 11":"Agriculture"}, axis = 1)
land.head()
```

Out[13]:

	LA name	Residential	Agriculture
0	Adur	4.226316	51.095201
1	Allerdale	0.282384	53.253852
2	Amber Valley	1.375413	67.644890
3	Arun	2.412087	55.505210
4	Ashfield	3.137516	44.549306


In [14]:

```
# merge with land usage
df = pd.merge(df, land, on='LA name')
df.head()
```

Out[14]:

	Accident_Index	Vehicle_Reference_x	Vehicle_Type	Towing_and_Articulation	Vehicle_Manoeuvr
0	2018010135259	1	9	0	-
1	2018010135259	2	9	0	-
2	2018010135309	1	9	0	1
3	2018010135309	2	11	0	1
4	2018010135320	1	1	-1	-

5 rows × 76 columns



In [15]:

```
# Final Set of Input Columns
df.columns
```

Out[15]:

```
Index(['Accident_Index', 'Vehicle_Reference_x', 'Vehicle_Type',
       'Towing_and_Articulation', 'Vehicle_Manoeuvre',
       'Vehicle_Location-Restricted_Lane', 'Junction_Location',
       'Skidding_and_Overturning', 'Hit_Object_in_Carriageway',
       'Vehicle_Leaving_Carriageway', 'Hit_Object_off_Carriageway',
       '1st_Point_of_Impact', 'Was_Vehicle_Left_Hand_Drive?',
       'Journey_Purpose_of_Driver', 'Sex_of_Driver', 'Age_of_Driver',
       'Age_Band_of_Driver', 'Engine_Capacity_(CC)', 'Propulsion_Code',
       'Age_of_Vehicle', 'Driver_IMD_Decile', 'Driver_Home_Area_Type',
       'Vehicle_IMD_Decile', 'Location_Easting_OSGR', 'Location_Northing_OSGR',
       'Longitude', 'Latitude', 'Police_Force', 'Accident_Severity',
       'Number_of_Vehicles', 'Number_of_Casualties', 'Date', 'Day_of_Week',
       'Time', 'Local_Authority_(District)', 'Local_Authority_(Highway)',
       '1st_Road_Class', '1st_Road_Number', 'Road_Type', 'Speed_limit',
       'Junction_Detail', 'Junction_Control', '2nd_Road_Class',
       '2nd_Road_Number', 'Pedestrian_Crossing-Human_Control',
       'Pedestrian_Crossing-Physical_Facilities', 'Light_Conditions',
       'Weather_Conditions', 'Road_Surface_Conditions',
       'Special_Conditions_at_Site', 'Carriageway_Hazards',
       'Urban_or_Rural_Area', 'Did_Police_Officer_Attend_Scene_of_Accident',
       'LSOA_of_Accident_Location', 'Vehicle_Reference_y',
       'Casualty_Reference', 'Casualty_Class', 'Sex_of_Casualty',
       'Age_of_Casualty', 'Age_Band_of_Casualty', 'Casualty_Severity',
       'Pedestrian_Location', 'Pedestrian_Movement', 'Car_Passenger',
       'Bus_or_Coach_Passenger', 'Pedestrian_Road_Maintenance_Worker',
       'Casualty_Type', 'Casualty_Home_Area_Type', 'Casualty_IMD_Decile',
       'code', 'LA name', 'Population', 'GDP', 'GDP_Growth', 'Residential',
       'Agriculture'],
      dtype='object')
```

In [16]:

```
# TODO - Fatalities (always) and older graphs
# TODO : Cleanup
# TODO : Missing
# TODO : PCA
# TODO : System Specs, GPU
# TODO : https://www3.cs.stonybrook.edu/~anshul/vis18\_poster.pdf
```

Descriptive Statistics


In [17]:

```
df = df.replace({-1:np.nan})
df.describe()
```

Out[17]:

	Vehicle_Reference_x	Vehicle_Type	Towing_and_Articulation	Vehicle_Manoeuvre	Vehicle_Restrict
count	232974.00000	232672.000000	231133.000000	228821.000000	2288
mean	1.65297	10.333560	0.031506	12.938358	
std	2.25571	10.898588	0.314701	6.185495	
min	1.00000	1.000000	0.000000	1.000000	
25%	1.00000	9.000000	0.000000	7.000000	
50%	1.00000	9.000000	0.000000	18.000000	
75%	2.00000	9.000000	0.000000	18.000000	
max	999.00000	98.000000	5.000000	18.000000	

8 rows × 70 columns



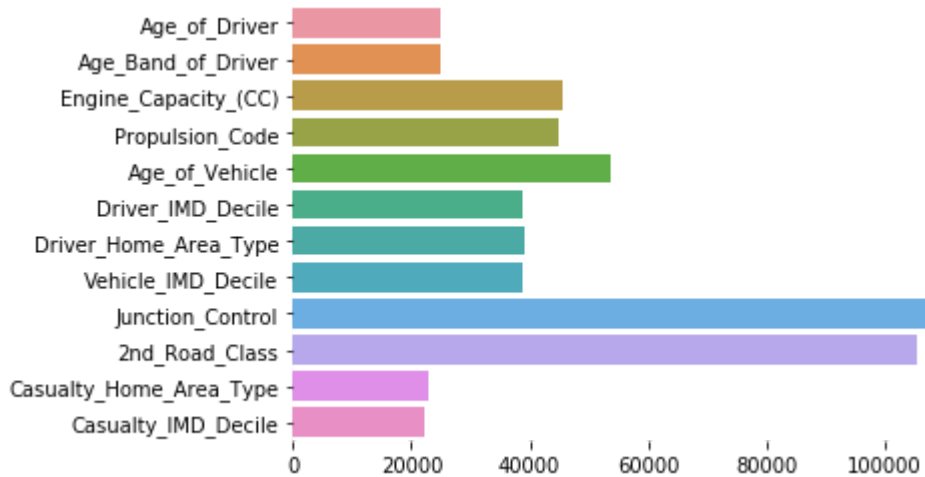
In [18]:

```
# Add New Variables
df['Month'] = pd.to_datetime(df.Date).dt.month
df['Hour'] = pd.to_datetime(df.Time).dt.hour
df['Older_Driver'] = (df['Age_of_Driver'] > 38)
```

In [19]:

```
missing = df.isna().sum()
missing = missing[missing > 5000]
sns.set_palette("muted")
ax = sns.barplot(y=missing.index, x=missing)
sns.despine(left=True, bottom=True)


# Drop Driver_IMD_Decile, Vehicle_IMD_Decile
df = df.drop(['Driver_IMD_Decile', 'Vehicle_IMD_Decile'], axis=1)
# Remove Unknown Gender
df = df[df['Sex_of_Driver'] != 3]
# Remove drivers with no age
df = df[df['Age_of_Driver'].notnull()]
```



EDA


In [20]:

```
# EDA of Gender
gender = df['Sex_of_Driver'].value_counts().rename({1.0:"Male", 2.0:"Female", 3.0:"Not Known"})
ax = sns.barplot(x=gender.index, y=gender)
ax.set(xlabel='', ylabel='Count')
sns.despine(left=True, bottom=True)
# Conclusion : More than 2x more Males than Females, need to be careful
# to make training data balanced between genders
```




In [21]:

```
males = df[df['Sex_of_Driver']==1]
females = df[df['Sex_of_Driver']==2]
fig, ax = plt.subplots(ncols=2, figsize=(12, 5), sharey=True)
boxplot = males['Age_of_Driver'].plot(kind='box', ax=ax[0], title='Male')
boxplot = females['Age_of_Driver'].plot(kind='box', ax=ax[1], title='Female')
```



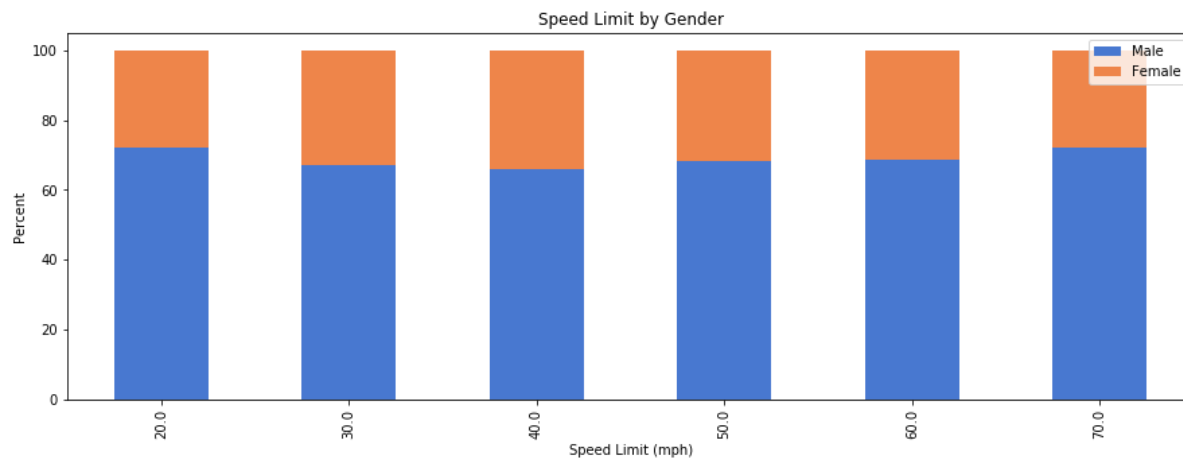
In [22]:

```
pvalue = stats.ks_2samp(males['Age_of_Driver'], females['Age_of_Driver']).pvalue
# Histogram of Age and Gender
fig, ax = plt.subplots()
sns.distplot(males['Age_of_Driver'], label='Male', hist=False, bins=range(1, 110, 10), ax=ax,
             kde=True)
sns.distplot(females['Age_of_Driver'], label='Female', hist=False, bins=range(1, 110, 10),
             ax=ax, kde=True)
ax.set(xlabel='Age', ylabel='Density')
ax.set_title("Age of Driver Distribution by Gender")
ax.text(80.0, 0.015, "KS Statistic\n\np-value =" + str(round(pvalue,2)))
plt.show()
```



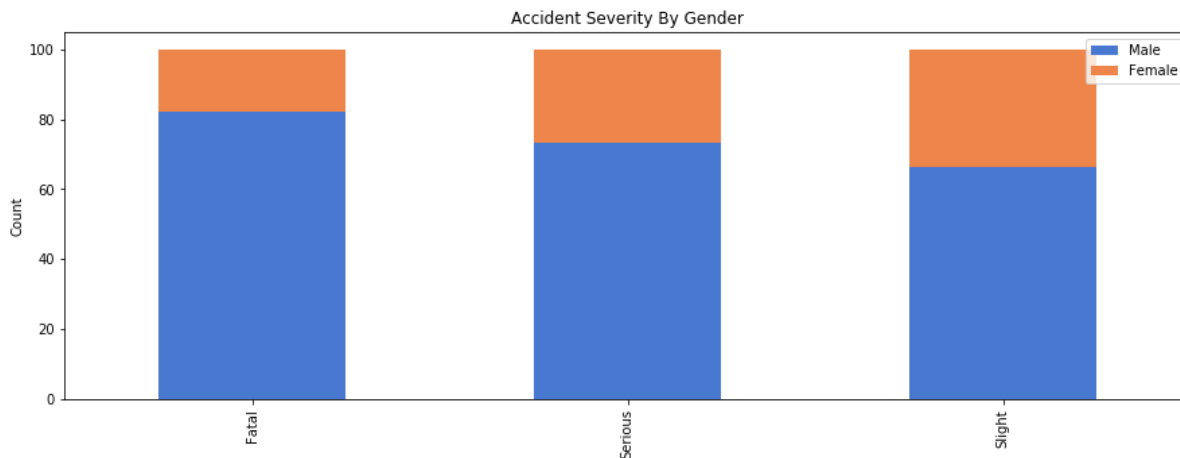
In [23]:

```
data = df.pivot_table(index='Speed_limit', columns='Sex_of_Driver', values='Road_Type', aggfunc='count')
data = data.dropna()
data = data.div(data.sum(1), axis=0) * 100
ax = data.plot(kind="bar", figsize=(15,5), stacked=True)
ax.set_xlabel("Speed Limit (mph)")
ax.set_ylabel("Percent")
ax.set_title("Speed Limit by Gender")
ax.legend(["Male", "Female"]);
```




In [24]:

```
data = df.pivot_table(index='Accident_Severity', columns='Sex_of_Driver', values='Road_Type', aggfunc='count')
data = data.dropna()
data = data.rename(index={1.0: "Fatal", 2.0: "Serious", 3.0: "Slight"})
data = data.div(data.sum(1), axis=0) * 100
ax = data.plot(kind="bar", figsize=(15,5), stacked=True)
ax.set_xlabel("")
ax.set_ylabel("Count")
ax.set_title("Accident Severity By Gender")
ax.legend(["Male", "Female"]);
```



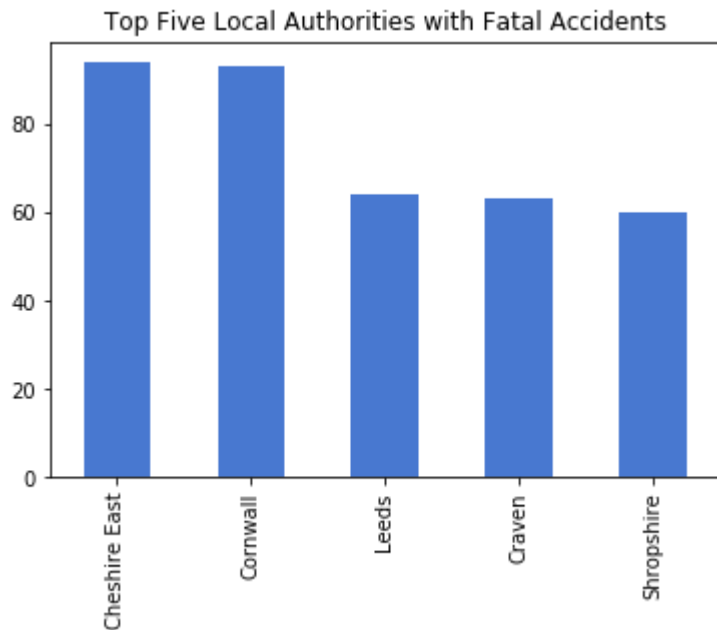
In [25]:

```
ax = df['LA name'].value_counts().iloc[0:5].plot(kind='bar')
x = ax.set_title("Top Five Local Authorities with Overall Accidents")
```



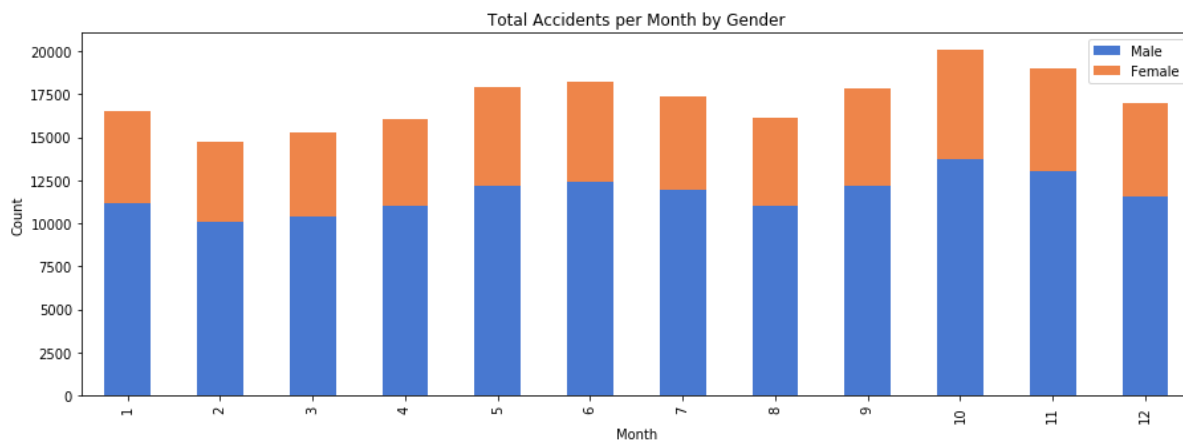
In [26]:

```
ax = df[df['Accident_Severity']==1]['LA name'].value_counts().iloc[0:5].plot(kind='bar')
x = ax.set_title("Top Five Local Authorities with Fatal Accidents")
```




In [27]:

```
data = df.pivot_table(index='Month', columns='Sex_of_Driver', values='Road_Type', aggfunc='count')
data = data.dropna()
ax = data.plot(kind="bar", figsize=(15,5), stacked=True)
ax.set_xlabel("Month")
ax.set_ylabel("Count")
ax.set_title("Total Accidents per Month by Gender")
ax.legend(["Male", "Female"]);
```



In [28]:

```
data = df[df['Accident_Severity']==1].pivot_table(index='Month', columns='Sex_of_Driver',
values='Road_Type', aggfunc='count')
data = data.dropna()
ax = data.plot(kind="bar", figsize=(15,5), stacked=True)
ax.set_xlabel("Month")
ax.set_ylabel("Count")
ax.set_title("Fatal Accidents per Month by Gender")
ax.legend(["Male", "Female"]);
```

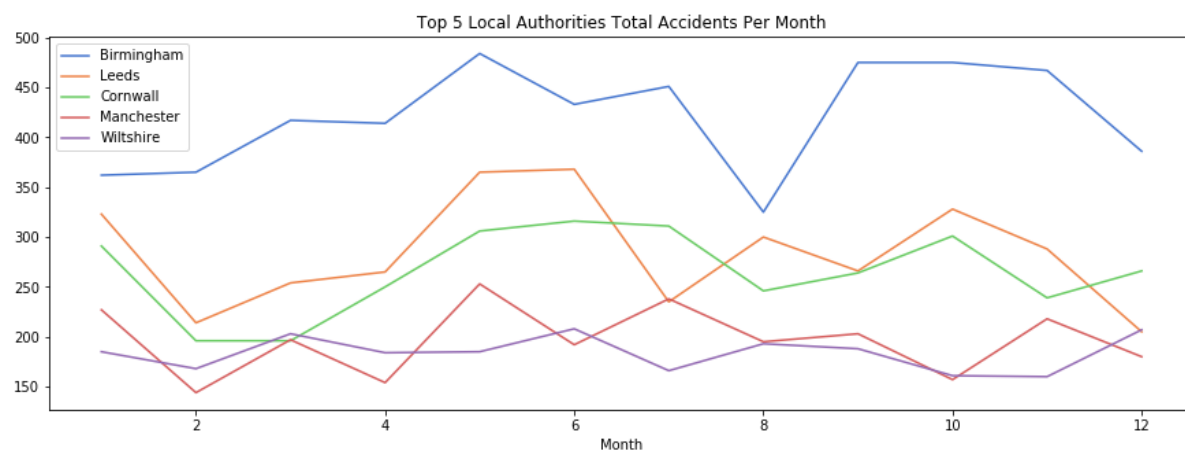


In [29]:

```
df[df["LA name"]== "Birmingham"].groupby("Month")["Road_Type"].count().plot(figsize=(15,5))
df[df["LA name"]== "Leeds"].groupby("Month")["Road_Type"].count().plot()
df[df["LA name"]== "Cornwall"].groupby("Month")["Road_Type"].count().plot()
df[df["LA name"]== "Manchester"].groupby("Month")["Road_Type"].count().plot()
df[df["LA name"]== "Wiltshire"].groupby("Month")["Road_Type"].count().plot()
plt.legend(["Birmingham", "Leeds", "Cornwall", "Manchester", "Wiltshire"])
plt.title("Top 5 Local Authorities Total Accidents Per Month", loc='center', pad=None)
```

Out[29]:

Text(0.5, 1.0, 'Top 5 Local Authorities Total Accidents Per Month')




In [30]:

```
df2 = df[df['Accident_Severity']==1]
df2[df2["LA name"]== "Cheshire East"].groupby("Month")["Road_Type"].count().plot(figsize=(15,5))
df2[df2["LA name"]== "Cornwall"].groupby("Month")["Road_Type"].count().plot()
df2[df2["LA name"]== "Leeds"].groupby("Month")["Road_Type"].count().plot()
df2[df2["LA name"]== "Craven"].groupby("Month")["Road_Type"].count().plot()
df2[df2["LA name"]== "Shropshire"].groupby("Month")["Road_Type"].count().plot()
plt.legend(["Cheshire East", "Cornwall", "Leeds", "Craven", "Shropshire"])
plt.title("Top 5 Local Authorities Fatal Accidents Per Month", loc='center', pad=None)
```

Out[30]:

Text(0.5, 1.0, 'Top 5 Local Authorities Fatal Accidents Per Month')

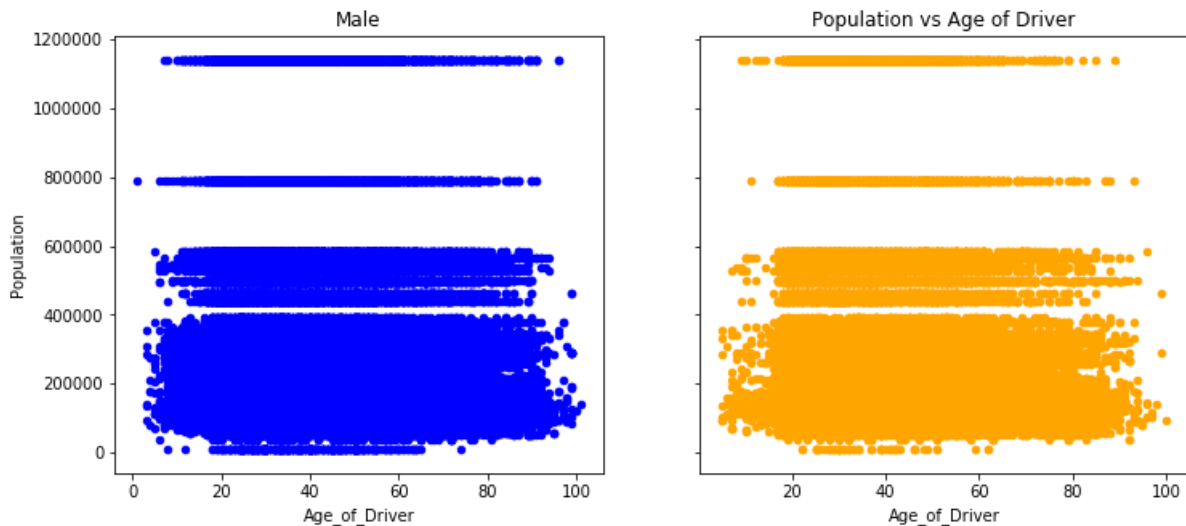


In [31]:

```
fig, ax = plt.subplots(ncols=2, figsize=(12, 5), sharey=True)
males.plot(kind='scatter', y='Population', x='Age_of_Driver', c='blue', ax=ax[0], title="Male")
females.plot(kind='scatter', y='Population', x='Age_of_Driver', c='orange', ax=ax[1], title="Female")
plt.title("Population vs Age of Driver", loc='center', pad=None)
```

Out[31]:

Text(0.5, 1.0, 'Population vs Age of Driver')




In [32]:

```
fig, ax = plt.subplots(ncols=2, figsize=(12, 5), sharey=True)
males.plot(kind='scatter', y='Residential', x='Age_of_Driver', c='blue', ax=ax[0], title=
"Male")
females.plot(kind='scatter', y='Residential', x='Age_of_Driver', c='orange', ax=ax[1], tit
le="Female")
```

Out[32]:

<matplotlib.axes._subplots.AxesSubplot at 0x18ed436af88>




In [33]:

```
fig, ax = plt.subplots(ncols=2, figsize=(12, 5), sharey=True)
males.plot(kind='scatter', y='Residential', x='Population', c='blue', ax=ax[0], title="Male")
females.plot(kind='scatter', y='Residential', x='Population', c='orange', ax=ax[1], title="Female")
```

Out[33]:

<matplotlib.axes._subplots.AxesSubplot at 0x18ed377f208>




In [34]:

```
df2 = df.pivot_table(index="Month", columns="Light_Conditions", values="Road_Type", aggfunc="count", fill_value=0)
df2 = df2.rename({1.0: "Daylight", 4.0: "Darkness - lights lit", 5.0: "Darkness - lights unlit", 6.0: "Darkness - no lighting", 7.0: "Darkness - lighting unknown"}, axis=1)
df2.plot(kind="bar", figsize=(15,5), stacked=True)
```

Out[34]:

<matplotlib.axes._subplots.AxesSubplot at 0x18ed37fd648>




In [35]:

```
df2 = df.pivot_table(index="Age_Band_of_Driver", columns="Light_Conditions", values="Road_Type", aggfunc="count", fill_value=0)
df2 = df2.rename({1.0: "0-10", 2.0: "6-10", 3.0: "11-15", 4.0: "16-20", 5.0: "21-25", 6.0: "26-35", 7.0: "36-45", 8.0: "45-55", 9.0: "56-65", 10.0: "66-75", 11.0: "Over 75"})
df2 = df2.rename({1.0: "Daylight", 4.0: "Darkness - lights lit", 5.0: "Darkness - lights unlit", 6.0: "Darkness - no lighting", 7.0: "Darkness - lighting unknown"}, axis=1)
df2.plot(kind="bar", figsize=(15,5), stacked=True)
```

Out[35]:

<matplotlib.axes._subplots.AxesSubplot at 0x18ed8b0e4c8>



In [36]:

```
# TODO : Time and Daytime etc, Road_Surface_Conditions etc
df = df[df.Towing_and_Articulation.notnull()]
df = df[df.Vehicle_Manoeuvre.notnull()]
df = df[df.Skidding_and_Overturning.notnull()]
df = df[df.Journey_Purpose_of_Driver.notnull()]
df = df[df.Casualty_Type.notnull()]
df = df[df.Sex_of_Casualty.notnull()]
df = df[df.Road_Surface_Conditions.notnull()]
df = df[df['GDP Growth'].notnull()]
df = df[df.Age_of_Casualty.notnull()]
df = df[df.Vehicle_Type.notnull()]
df = df[df.Hour.notnull()]

print("Number of datapoints after removal of missing data = ", df.shape[0])
```

Number of datapoints after removal of missing data = 198414

Pipeline

In [37]:

```
# z-scale for quantative data
def z_scale(x):
    return (x-x.mean())/x.std()
def nop(x):
    return x

# TODO : Binary : Is Weekend?

# column transformer
column_transformer = ColumnTransformer(
    transformers = [
        ('Towing_and_Articulation', OneHotEncoder(categories='auto'), ["Towing_and_Articulation"]),
        ('Vehicle_Type', OneHotEncoder(categories='auto'), ["Vehicle_Type"]),
        ('Vehicle_Manoeuvre', OneHotEncoder(categories='auto'), ["Vehicle_Manoeuvre"]),
        ('Skidding_and_Overturning', OneHotEncoder(categories='auto'), ["Skidding_and_Overturning"]),
        ('Journey_Purpose_of_Driver', OneHotEncoder(categories='auto'), ["Journey_Purpose_of_Driver"]),
        ('LA name', OneHotEncoder(categories='auto'), ["LA name"]),
        ('Hour', OneHotEncoder(categories='auto'), ["Hour"]),
        ('Road_Type', OneHotEncoder(categories='auto'), ["Road_Type"]),
        ('Month', OrdinalEncoder(), ["Month"]),
        ('Sex_of_Casualty', OneHotEncoder(categories='auto'), ["Sex_of_Casualty"]),
        ('Sex_of_Driver', OneHotEncoder(categories='auto'), ["Sex_of_Driver"]),
        ('Number_of_Casualties', MinMaxScaler(), ["Number_of_Casualties"]),
        ('Weather_Conditions', OneHotEncoder(categories='auto'), ["Weather_Conditions"]),
        ('Road_Surface_Conditions', OneHotEncoder(categories='auto'), ["Road_Surface_Conditions"]),
        ('Light_Conditions', OneHotEncoder(categories='auto'), ["Light_Conditions"]),
        ('Urban_or_Rural_Area', OneHotEncoder(categories='auto'), ["Urban_or_Rural_Area"]),
        ('Population', MinMaxScaler(), ["Population"]),
        ('GDP', FunctionTransformer(func=z_scale, validate=False), ["GDP"]),
        ('GDP Growth', FunctionTransformer(func=nop, validate=False), ["GDP Growth"]),
        ('Age_of_Casualty', FunctionTransformer(func=nop, validate=False), ["Age_of_Casualty"]),
    ],
    ('Agriculture', MinMaxScaler(), ["Agriculture"]),
    ('Residential', MinMaxScaler(), ["Residential"]),
    #('Age_of_Driver', FunctionTransformer(func=nop, validate=False), ["Age_of_Driver"]),
    ],
    remainder = 'drop')

def show_confusion_result(title, x, y):
    start = time.time()
    y_pred = pipeline.predict(x)
    end = time.time()
    # create confusion matrix
    cf_matrix = confusion_matrix(y_true=y, y_pred=y_pred)
    TN, FP, FN, TP = cf_matrix.ravel()
    precision = TP/(TP+FP)
    recall = TP/(TP+FN)
    F1 = 2*(precision*recall)/(precision+recall)
    accuracy = (TP+TN)/(TN + FP + FN + TP)
    # Work Cited: https://medium.com/@dtuk81/confusion-matrix-visualization-fc31e3f30fea
    group_names = ['True Neg', 'False Pos', 'False Neg', 'True Pos']
```



```

group_counts = ["{0:0.0f}".format(value) for value in cf_matrix.flatten()]
group_percentages = ["{0:.2%}".format(value) for value in cf_matrix.flatten()/np.sum(c
f_matrix)]
labels = [f"{v1}\n\n{v2}\n\n{v3}" for v1, v2, v3 in zip(group_names,group_counts,group
_percentages)]
labels = np.asarray(labels).reshape(2,2)
fig, ax = plt.subplots()
sns.heatmap(cf_matrix, annot=labels, fmt='', cmap='Blues', ax=ax)
ax.set_ylim([0,2])
ax.text(2.7, 1.4, 'Accuracy ' + str(round(accuracy*100,2)) + '%', c='blue')
ax.text(2.7, 1.2, 'Precision ' + str(round(precision*100,2)) + '%', c='blue')
ax.text(2.7, 1.0, 'Recall ' + str(round(recall*100,2)) + '%', c='blue' )
ax.text(2.7, 0.8, 'F1 ' + str(round(F1*100,2)) + '%', c='blue' )
x=plt.title(title)
return [accuracy, precision, recall, F1, end-start]

def make_summary_table(result):
    results_summary = (pd.DataFrame(np.array(result), columns=['Accuracy %', 'Precision %'
, 'Recall %', 'F1 %', 'Prediction Time (s)'])*100.0).round(2)
    results_summary['Settings'] = ['Default', 'Default', 'Default', 'Optimized', 'Optimize
d', 'Optimized']
    results_summary['Data'] = ['Train', 'Validate', 'Test', 'Train', 'Validate', 'Test']
    results_summary = results_summary[['Settings', 'Data', 'Accuracy %', 'Precision %', 'R
ecall %', 'F1 %', 'Prediction Time (s)']]
    return results_summary

fit_time = []

```

Train (60), Validate (20), Test (20) Split

In [38]:

```

label = 'Older_Driver'
# split it into training, validation, and test splits, with a 60/20/20%
x_train_validate, x_test, y_train_validate, y_test = train_test_split(df.drop(label, axis=
1), df[label], test_size=0.2, random_state=42)
# now split train_validate into train (75%) and validate (25%)
x_train, x_validate, y_train, y_validate = train_test_split(x_train_validate, y_train_vali
date, test_size=0.25, random_state=42)
# set number of folds for validation:
nfolds = 5

```

Decision Tree Classifier (default)

In [39]:

```

decisiontree = DecisionTreeClassifier()
pipeline = Pipeline([
    ('preprocessing', column_transformer),
    ('decisiontree', decisiontree)
], verbose=verbose)

```

In [40]:




```
# Train
pipeline.fit(x_train, y_train)
# Accuracy
train_score = accuracy_score(y_train, pipeline.predict(x_train), normalize=True, sample_weight=None)
test_score = accuracy_score(y_test, pipeline.predict(x_test), normalize=True, sample_weight=None)
train_score, test_score
```

Out[40]:

```
(0.9834352530071904, 0.6800141118363027)
```

In [41]:

```
result = []
result.append(show_confusion_result('Confusion Matrix : Train (Default)', x_train, y_train
))
result.append(show_confusion_result('Confusion Matrix : Validate (Default)', x_validate, y
_validate))
result.append(show_confusion_result('Confusion Matrix : Test (Default)', x_test, y_test))
```



Decision Tree Classifier Hyper-Parameter Tuning

In [42]:

```
criterion = ['gini', 'entropy']
splitter = ['best', 'random']
max_depth = [6,10,20]
parameters = dict(decisiontree__criterion=criterion,
                  decisiontree__max_depth=max_depth,
                  decisiontree__splitter = splitter)
clf = GridSearchCV(pipeline, parameters, cv=nfolds)
# Fit the grid search
clf.fit(x_validate, y_validate)
print('Best criterion:', clf.best_estimator_.get_params()['decisiontree__criterion'])
print('Best max_depth:', clf.best_estimator_.get_params()['decisiontree__max_depth'])
print('Best splitter:', clf.best_estimator_.get_params()['decisiontree__splitter'])
```

Best criterion: gini
Best max_depth: 6
Best splitter: best

In [43]:




```
decisiontree = clf.best_estimator_.get_params()['decisiontree']
pipeline = Pipeline([
    ('preprocessing', column_transformer),
    ('decisiontree', decisiontree)
], verbose=verbose)
# Train
start = time.time();
pipeline.fit(x_train, y_train);
fit_time.append(time.time()-start);

# Accuracy
train_score = accuracy_score(y_train, pipeline.predict(x_train), normalize=True, sample_weight=None)
print("Training Data Accuracy =", train_score, " Test Data Accuracy =", test_score)
```

Training Data Accuracy = 0.7201968953699348 Test Data Accuracy = 0.6800141118363027

In [44]:

```
result.append(show_confusion_result('Confusion Matrix : Test (Optimized)', x_train, y_train))
result.append(show_confusion_result('Confusion Matrix : Test (Optimized)', x_validate, y_validate))
result.append(show_confusion_result('Confusion Matrix : Test (Optimized)', x_test, y_test))
```



Decision Tree Final Results

In [45]:

```
decisiontree_df = make_summary_table(result)
decisiontree_df
```

Out[45]:

	Settings	Data	Accuracy %	Precision %	Recall %	F1 %	Prediction Time (s)
0	Default	Train	98.34	99.69	97.01	98.33	37.19
1	Default	Validate	67.69	68.75	66.39	67.55	15.62
2	Default	Test	68.00	68.47	67.17	67.81	17.89
3	Optimized	Train	72.02	74.65	67.25	70.75	37.84
4	Optimized	Validate	71.71	74.77	66.61	70.45	14.69
5	Optimized	Test	71.66	74.18	66.76	70.28	15.62

In []:

Random Forest Classifier (default)

In [46]:

```
randomforest = RandomForestClassifier(n_estimators=5,max_depth=5)
pipeline = Pipeline([
    ('preprocessing', column_transformer),
    ('randomforest', randomforest)
], verbose=verbose)
```

In [47]:




```
# Train
pipeline.fit(x_train, y_train)
# Accuracy
train_score = accuracy_score(y_train, pipeline.predict(x_train), normalize=True, sample_weight=None)
test_score = accuracy_score(y_test, pipeline.predict(x_test), normalize=True, sample_weight=None)
train_score, test_score
```

Out[47]:

```
(0.5902073113366038, 0.5864475972078724)
```

In [48]:

```
result = []
result.append(show_confusion_result('Confusion Matrix : Train (Default)', x_train, y_train
))
result.append(show_confusion_result('Confusion Matrix : Validate (Default)', x_validate, y
_validate))
result.append(show_confusion_result('Confusion Matrix : Test (Default)', x_test, y_test))
```

Random Forest Classifier Hyper-Parameter Tuning

In [49]:

```
n_estimators = [5,20]
max_depth = [5,30]
criterion = ['entropy','gini']

parameters = dict(randomforest__n_estimators=n_estimators,
                  randomforest__max_depth=max_depth,
                  randomforest__criterion = criterion)
clf = GridSearchCV(pipeline, parameters, cv=nfolds)
# Fit the grid search
clf.fit(x_validate, y_validate)
print('Best n_estimators:', clf.best_estimator_.get_params()['randomforest__n_estimators'])
print('Best max_depth:', clf.best_estimator_.get_params()['randomforest__max_depth'])
print('Best criterion:', clf.best_estimator_.get_params()['randomforest__criterion'])
```

Best n_estimators: 20
Best max_depth: 30
Best criterion: gini

In [50]:




```
randomforest = clf.best_estimator_.get_params()['randomforest']
pipeline = Pipeline([
    ('preprocessing', column_transformer),
    ('randomforest', randomforest)
], verbose=verbose)
# Train
start = time.time();
pipeline.fit(x_train, y_train);
fit_time.append(time.time()-start);

# Accuracy
test_score = accuracy_score(y_test, pipeline.predict(x_test), normalize=True, sample_weight=None)
test_score = accuracy_score(y_test, pipeline.predict(x_test), normalize=True, sample_weight=None)
print("Training Data Accuracy =", train_score, " Test Data Accuracy =", test_score)
```

Training Data Accuracy = 0.5902073113366038 Test Data Accuracy = 0.745986946
55142

In [51]:

```
result.append(show_confusion_result('Confusion Matrix : Test (Optimized)', x_train, y_train))
result.append(show_confusion_result('Confusion Matrix : Test (Optimized)', x_validate, y_validate))
result.append(show_confusion_result('Confusion Matrix : Test (Optimized)', x_test, y_test))
```



Random Forest Final Results

In [52]:

```
randomforest_df = make_summary_table(result)
randomforest_df
```

Out[52]:

	Settings	Data	Accuracy %	Precision %	Recall %	F1 %	Prediction Time (s)
0	Default	Train	59.02	56.78	77.77	65.64	40.74
1	Default	Validate	59.04	57.02	77.66	65.76	15.56
2	Default	Test	58.64	56.38	77.73	65.36	17.19
3	Optimized	Train	86.24	86.22	86.49	86.36	97.95
4	Optimized	Validate	74.58	75.27	74.19	74.72	38.36
5	Optimized	Test	74.60	74.88	74.32	74.60	37.50

In []:

SVC Classifier (default)

In [53]:

```
svc = SVC(gamma='auto', max_iter=10)
pipeline = Pipeline([
    ('preprocessing', column_transformer),
    ('svc', svc)
],
verbose=verbose)
```

In [54]:




```
# Train
pipeline.fit(x_train, y_train)
# Accuracy
train_score = accuracy_score(y_train, pipeline.predict(x_train), normalize=True, sample_weight=None)
test_score = accuracy_score(y_test, pipeline.predict(x_test), normalize=True, sample_weight=None)
train_score, test_score
```

Out[54]:

```
(0.6215476110476447, 0.6206436005342338)
```

In [55]:

```
result = []  
result.append(show_confusion_result('Confusion Matrix : Train (Default)', x_train, y_train  
)  
)  
result.append(show_confusion_result('Confusion Matrix : Validate (Default)', x_validate, y  
_validate))  
result.append(show_confusion_result('Confusion Matrix : Test (Default)', x_test, y_test))
```



SVC Classifier Results Hyper-Parameter Tuning

In [56]:

```
C = [(10 ** i) for i in range(-1,5)]

parameters = dict(svc__C=C)
clf = GridSearchCV(pipeline, parameters, cv=nfolds)
# Fit the grid search
clf.fit(x_validate, y_validate)
print('Best C:', clf.best_estimator_.get_params()['svc__C'])
```

Best C: 0.1

In [57]:



```
svc = clf.best_estimator_.get_params()['svc']
pipeline = Pipeline([
    ('preprocessing', column_transformer),
    ('svc', svc)
])
# Train
start = time.time();
pipeline.fit(x_train, y_train);
fit_time.append(time.time()-start);


# Accuracy
train_score = accuracy_score(y_train, pipeline.predict(x_train), normalize=True, sample_weight=None)
test_score = accuracy_score(y_test, pipeline.predict(x_test), normalize=True, sample_weight=None)
print("Training Data Accuracy =", train_score, " Test Data Accuracy =", test_score)
```

Training Data Accuracy = 0.6215476110476447 Test Data Accuracy = 0.6206436005342338

In [58]:

```
result.append(show_confusion_result('Confusion Matrix : Test (Optimized)', x_train, y_train))
result.append(show_confusion_result('Confusion Matrix : Test (Optimized)', x_validate, y_validate))
result.append(show_confusion_result('Confusion Matrix : Test (Optimized)', x_test, y_test))
```



SVC Final Results

In [59]:

```
svc_df = make_summary_table(result)
svc_df
```

Out[59]:

	Settings	Data	Accuracy %	Precision %	Recall %	F1 %	Prediction Time (s)
0	Default	Train	62.15	68.22	46.45	55.27	64.16
1	Default	Validate	62.14	68.57	46.61	55.49	23.66
2	Default	Test	62.06	67.77	46.55	55.19	24.91
3	Optimized	Train	62.15	68.22	46.45	55.27	54.91
4	Optimized	Validate	62.14	68.57	46.61	55.49	22.26
5	Optimized	Test	62.06	67.77	46.55	55.19	25.12

In []:

Gradient Boost Classifier (default)

In [60]:

```
gboost = GradientBoostingClassifier()
pipeline = Pipeline([
    ('preprocessing', column_transformer),
    ('gboost', gboost)
], verbose=verbose)
```

In [61]:

```
# Train
pipeline.fit(x_train, y_train)
# Accuracy
train_score = accuracy_score(y_train, pipeline.predict(x_train), normalize=True, sample_weight=None)
test_score = accuracy_score(y_test, pipeline.predict(x_test), normalize=True, sample_weight=None)
train_score, test_score
```


Out[61]:

```
(0.7292268664740272, 0.725096388881889)
```


In [62]:

```
result = []
result.append(show_confusion_result('Confusion Matrix : Train (Default)', x_train, y_train
))
result.append(show_confusion_result('Confusion Matrix : Validate (Default)', x_validate, y
_validate))
result.append(show_confusion_result('Confusion Matrix : Test (Default)', x_test, y_test))
```


Confusion Matrix : Train (Default)



Confusion Matrix : Validate (Default)



Confusion Matrix : Test (Default)



Gradient Boost Hyper-Parameter Tuning

In [63]:

```
loss = ['deviance', 'exponential']
n_estimators = [50,100,200]

parameters = dict(gboost__n_estimators=n_estimators,
                  gboost__loss = loss)
clf = GridSearchCV(pipeline, parameters, cv=nfolds)
# Fit the grid search
clf.fit(x_validate, y_validate)
print('Best n_estimators:', clf.best_estimator_.get_params()['gboost__n_estimators'])
print('Best loss:', clf.best_estimator_.get_params()['gboost__loss'])
```

Best n_estimators: 200

Best loss: deviance

In [64]:



```
gboost = clf.best_estimator_.get_params()['gboost']
pipeline = Pipeline([
    ('preprocessing', column_transformer),
    ('gboost', gboost)
], verbose=verbose)
# Train
start = time.time();
pipeline.fit(x_train, y_train);
fit_time.append(time.time()-start);


# Accuracy
train_score = accuracy_score(y_train, pipeline.predict(x_train), normalize=True, sample_weight=None)
test_score = accuracy_score(y_test, pipeline.predict(x_test), normalize=True, sample_weight=None)
print("Training Data Accuracy =", train_score, " Test Data Accuracy =", test_score)
```

Training Data Accuracy = 0.7314276594314898 Test Data Accuracy = 0.7269107678350931

In [65]:

```
result.append(show_confusion_result('Confusion Matrix : Test (Optimized)', x_train, y_train))
result.append(show_confusion_result('Confusion Matrix : Test (Optimized)', x_validate, y_validate))
result.append(show_confusion_result('Confusion Matrix : Test (Optimized)', x_test, y_test))
```





Gradient Boost Final Results

In [66]:

```
gboost_df = make_summary_table(result)
gboost_df
```

Out[66]:

	Settings	Data	Accuracy %	Precision %	Recall %	F1 %	Prediction Time (s)
0	Default	Train	72.92	73.76	71.71	72.72	53.27
1	Default	Validate	72.46	73.63	71.08	72.33	21.87
2	Default	Test	72.51	73.30	71.14	72.20	18.24
3	Optimized	Train	73.14	73.55	72.83	73.19	91.03
4	Optimized	Validate	72.68	73.43	72.17	72.79	31.50
5	Optimized	Test	72.69	73.06	72.22	72.64	26.56

In []:

Bagging Classifier (default)

In [67]:

```
bagging = BaggingClassifier()
pipeline = Pipeline([
    ('preprocessing', column_transformer),
    ('bagging', bagging)
], verbose=verbose)
```

In [68]:

```
# Train
pipeline.fit(x_train, y_train)
# Accuracy
train_score = accuracy_score(y_train, pipeline.predict(x_train), normalize=True, sample_weight=None)
test_score = accuracy_score(y_test, pipeline.predict(x_test), normalize=True, sample_weight=None)
train_score, test_score
```


Out[68]:

```
(0.9695500974396882, 0.7131265277322784)
```


In [69]:

```
result = []
result.append(show_confusion_result('Confusion Matrix : Train (Default)', x_train, y_train
))
result.append(show_confusion_result('Confusion Matrix : Validate (Default)', x_validate, y
_validate))
result.append(show_confusion_result('Confusion Matrix : Test (Default)', x_test, y_test))
```


Confusion Matrix : Train (Default)



Confusion Matrix : Validate (Default)



Confusion Matrix : Test (Default)



Bagging Classifier Hyper-Parameter Tuning

In [70]:

```
max_samples = [0.5, 1.0, 1.5]
max_features = [0.5, 1.0, 1.5]
parameters = dict(bagging__max_features = max_features,
                  bagging__max_samples = max_samples)
clf = GridSearchCV(pipeline, parameters, cv=nfolds)
# Fit the grid search
clf.fit(x_validate, y_validate)
print('Best max_features:', clf.best_estimator_.get_params()['bagging__max_features'])
print('Best _max_samples:', clf.best_estimator_.get_params()['bagging__max_samples'])
```

Best max_features: 1.0

Best _max_samples: 0.5

In [71]:



```
bagging = clf.best_estimator_.get_params()['bagging']
pipeline = Pipeline([
    ('preprocessing', column_transformer),
    ('bagging', bagging)
], verbose=verbose)
# Train
start = time.time();
pipeline.fit(x_train, y_train);
fit_time.append(time.time()-start);


# Accuracy
train_score = accuracy_score(y_train, pipeline.predict(x_train), normalize=True, sample_weight=None)
test_score = accuracy_score(y_test, pipeline.predict(x_test), normalize=True, sample_weight=None)
print("Training Data Accuracy =", train_score, " Test Data Accuracy =", test_score)
```

Training Data Accuracy = 0.9074490961628923 Test Data Accuracy = 0.7189224604994582

In [72]:

```
result.append(show_confusion_result('Confusion Matrix : Test (Optimized)', x_train, y_train))
result.append(show_confusion_result('Confusion Matrix : Test (Optimized)', x_validate, y_validate))
result.append(show_confusion_result('Confusion Matrix : Test (Optimized)', x_test, y_test))
```





Bagging Final Results

In [73]:

```
bagging_df = make_summary_table(result)
bagging_df
```

Out[73]:

	Settings	Data	Accuracy %	Precision %	Recall %	F1 %	Prediction Time (s)
0	Default	Train	96.96	97.51	96.41	96.96	107.39
1	Default	Validate	71.28	72.73	69.26	70.95	39.06
2	Default	Test	71.31	72.23	69.59	70.89	42.18
3	Optimized	Train	90.74	92.49	88.82	90.62	103.89
4	Optimized	Validate	71.77	73.36	69.49	71.37	37.63
5	Optimized	Test	71.89	73.11	69.59	71.31	34.34

In []:

Naive Bayes (default)

In [74]:

```
nbayes = BernoulliNB()
pipeline = Pipeline([
    ('preprocessing', column_transformer),
    ('nbayes', nbayes)
], verbose=verbose)
```


In [75]:

```
# Train
pipeline.fit(x_train, y_train)
# Accuracy
train_score = accuracy_score(y_train, pipeline.predict(x_train), normalize=True, sample_weight=None)
test_score = accuracy_score(y_test, pipeline.predict(x_test), normalize=True, sample_weight=None)
train_score, test_score
```


Out[75]:

```
(0.5884517169545057, 0.5843056220547842)
```


In [76]:

```
result = []
result.append(show_confusion_result('Confusion Matrix : Train (Default)', x_train, y_train
))
result.append(show_confusion_result('Confusion Matrix : Validate (Default)', x_validate, y
_validate))
result.append(show_confusion_result('Confusion Matrix : Test (Default)', x_test, y_test))
```


Confusion Matrix : Train (Default)



Confusion Matrix : Validate (Default)



Confusion Matrix : Test (Default)



Naive Bayes Hyper-Parameter Tuning

In [77]:

```
alpha = [0.5, 1.0, 1.5]
binarize = [0.5, 1.0, 1.5]
parameters = dict(nbayes__alpha = alpha,
                  nbayes__binarize = binarize)
clf = GridSearchCV(pipeline, parameters, cv=nfolds)
# Fit the grid search
clf.fit(x_validate, y_validate)
print('Best alpha:', clf.best_estimator_.get_params()['nbayes__alpha'])
print('Best binarize:', clf.best_estimator_.get_params()['nbayes__binarize'])
```

Best alpha: 1.0

Best binarize: 0.5

In [78]:



```
nbayes = clf.best_estimator_.get_params()['nbayes']
pipeline = Pipeline([
    ('preprocessing', column_transformer),
    ('nbayes', nbayes)
], verbose=verbose)
# Train
start = time.time();
pipeline.fit(x_train, y_train);
fit_time.append(time.time()-start);


# Accuracy
train_score = accuracy_score(y_train, pipeline.predict(x_train), normalize=True, sample_weight=None)
test_score = accuracy_score(y_test, pipeline.predict(x_test), normalize=True, sample_weight=None)
print("Training Data Accuracy =", train_score, " Test Data Accuracy =", test_score)
```

Training Data Accuracy = 0.5885189167394664 Test Data Accuracy = 0.5841796235163672

In [79]:

```
result.append(show_confusion_result('Confusion Matrix : Test (Optimized)', x_train, y_train))
result.append(show_confusion_result('Confusion Matrix : Test (Optimized)', x_validate, y_validate))
result.append(show_confusion_result('Confusion Matrix : Test (Optimized)', x_test, y_test))
```





Naive Bayes Final Results

In [80]:

```
nbayes_df = make_summary_table(result)
nbayes_df
```

Out[80]:

	Settings	Data	Accuracy %	Precision %	Recall %	F1 %	Prediction Time (s)
0	Default	Train	58.85	58.29	64.12	61.07	39.56
1	Default	Validate	58.75	58.47	64.04	61.13	14.06
2	Default	Test	58.43	57.76	63.91	60.68	14.06
3	Optimized	Train	58.85	58.47	63.00	60.65	43.24
4	Optimized	Validate	58.76	58.64	63.02	60.75	17.19
5	Optimized	Test	58.42	57.92	62.71	60.22	17.18

In []:

Nearest Centroid (default)

In [81]:

```
nccentroid = NearestCentroid()
pipeline = Pipeline([
    ('preprocessing', column_transformer),
    ('nccentroid', nccentroid)
], verbose=verbose)
```

In [82]:




```
# Train
pipeline.fit(x_train, y_train)
# Accuracy
train_score = accuracy_score(y_train, pipeline.predict(x_train), normalize=True, sample_weight=None)
test_score = accuracy_score(y_test, pipeline.predict(x_test), normalize=True, sample_weight=None)
train_score, test_score
```

Out[82]:

```
(0.7181473019286339, 0.7151425043469496)
```


In [83]:

```
result = []
result.append(show_confusion_result('Confusion Matrix : Train (Default)', x_train, y_train
))
result.append(show_confusion_result('Confusion Matrix : Validate (Default)', x_validate, y
_validate))
result.append(show_confusion_result('Confusion Matrix : Test (Default)', x_test, y_test))
```



Nearest Centroid Hyper-Parameter Tuning

In [84]:

```
metric = ['euclidean', 'manhattan']
parameters = dict(ncentroid__metric = metric)
clf = GridSearchCV(pipeline, parameters, cv=nfolds)
# Fit the grid search
clf.fit(x_validate, y_validate)
print('Best metric:', clf.best_estimator_.get_params()['ncentroid__metric'])
```

Best metric: euclidean

In [85]:



```
ncentroid = clf.best_estimator_.get_params()['ncentroid']
pipeline = Pipeline([
    ('preprocessing', column_transformer),
    ('ncentroid', ncentroid)
], verbose=verbose)
# Train
start = time.time();
pipeline.fit(x_train, y_train);
fit_time.append(time.time()-start);


# Accuracy
train_score = accuracy_score(y_train, pipeline.predict(x_train), normalize=True, sample_weight=None)
test_score = accuracy_score(y_test, pipeline.predict(x_test), normalize=True, sample_weight=None)
print("Training Data Accuracy =", train_score, " Test Data Accuracy =", test_score)
```

Training Data Accuracy = 0.7181473019286339 Test Data Accuracy = 0.7151425043469496

In [86]:

```
result.append(show_confusion_result('Confusion Matrix : Test (Optimized)', x_train, y_train))
result.append(show_confusion_result('Confusion Matrix : Test (Optimized)', x_validate, y_validate))
result.append(show_confusion_result('Confusion Matrix : Test (Optimized)', x_test, y_test))
```





Nearest Centroid Final Results

In [87]:

```
ncentroid_df = make_summary_table(result)
ncentroid_df
```

Out[87]:

	Settings	Data	Accuracy %	Precision %	Recall %	F1 %	Prediction Time (s)
0	Default	Train	71.81	74.58	66.75	70.45	36.77
1	Default	Validate	71.51	74.72	66.11	70.15	15.62
2	Default	Test	71.51	74.17	66.34	70.04	14.06
3	Optimized	Train	71.81	74.58	66.75	70.45	38.11
4	Optimized	Validate	71.51	74.72	66.11	70.15	15.74
5	Optimized	Test	71.51	74.17	66.34	70.04	15.62

In []:

Multi Layer Perceptron (default)

In [88]:

```
mlp = MLPClassifier()
pipeline = Pipeline([
    ('preprocessing', column_transformer),
    ('mlp', mlp)
], verbose=verbose)
```

In [89]:




```
# Train
pipeline.fit(x_train, y_train)
# Accuracy
train_score = accuracy_score(y_train, pipeline.predict(x_train), normalize=True, sample_weight=None)
test_score = accuracy_score(y_test, pipeline.predict(x_test), normalize=True, sample_weight=None)
train_score, test_score
```

Out[89]:

```
(0.8609888448356965, 0.6982587001990777)
```

In [90]:

```
result = []
result.append(show_confusion_result('Confusion Matrix : Train (Default)', x_train, y_train
))
result.append(show_confusion_result('Confusion Matrix : Validate (Default)', x_validate, y
_validate))
result.append(show_confusion_result('Confusion Matrix : Test (Default)', x_test, y_test))
```

Multi Layer Perceptron Hyper-Parameter Tuning

In [91]:

```
nfolds = 2 # too slow for too many folds
activation = ['logistic', 'tanh', 'relu']
solver = ['lbfgs', 'adam']

parameters = dict(mlp__activation = activation,
                  mlp__solver = solver)
clf = GridSearchCV(pipeline, parameters, cv=nfolds)
# Fit the grid search
clf.fit(x_validate, y_validate)
print('Best activation:', clf.best_estimator_.get_params()['mlp__activation'])
print('Best solver:', clf.best_estimator_.get_params()['mlp__solver'])
```

Best activation: logistic

Best solver: lbfgs

In [92]:




```
mlp = clf.best_estimator_.get_params()['mlp']
pipeline = Pipeline([
    ('preprocessing', column_transformer),
    ('mlp', mlp)
], verbose=verbose)
# Train
start = time.time();
pipeline.fit(x_train, y_train);
fit_time.append(time.time()-start);

# Accuracy
train_score = accuracy_score(y_train, pipeline.predict(x_train), normalize=True, sample_weight=None)
test_score = accuracy_score(y_test, pipeline.predict(x_test), normalize=True, sample_weight=None)
print("Training Data Accuracy =", train_score, " Test Data Accuracy =", test_score)
```

Training Data Accuracy = 0.7179289026275116 Test Data Accuracy = 0.7131517274399617

In [93]:

```
result.append(show_confusion_result('Confusion Matrix : Test (Optimized)', x_train, y_train))
result.append(show_confusion_result('Confusion Matrix : Test (Optimized)', x_validate, y_validate))
result.append(show_confusion_result('Confusion Matrix : Test (Optimized)', x_test, y_test))
```



Multi Layer Perceptron Final Results

In [94]:

```
mlp_df = make_summary_table(result)
mlp_df
```

Out[94]:

	Settings	Data	Accuracy %	Precision %	Recall %	F1 %	Prediction Time (s)
0	Default	Train	86.10	85.87	86.63	86.25	143.55
1	Default	Validate	69.41	69.93	69.45	69.69	48.57
2	Default	Test	69.83	69.86	70.13	70.00	48.54
3	Optimized	Train	71.79	71.28	73.61	72.43	51.69
4	Optimized	Validate	71.11	70.87	72.93	71.89	17.19
5	Optimized	Test	71.32	70.72	73.12	71.90	18.75

Model Final Results

In [95]:

```
summary_df = decisiontree_df.assign(Model="Decision Tree").assign(Fit_Time=fit_time[0])
summary_df = summary_df.append(randomforest_df.assign(Model="Random Forest").assign(Fit_Time=fit_time[1]))
summary_df = summary_df.append(svc_df.assign(Model="SVC").assign(Fit_Time=fit_time[2]))
summary_df = summary_df.append(gboost_df.assign(Model="Gradient Boosting").assign(Fit_Time=fit_time[3]))
summary_df = summary_df.append(bagging_df.assign(Model="Bagging").assign(Fit_Time=fit_time[4]))
summary_df = summary_df.append(nbayes_df.assign(Model="Naive Bayes (Bernoulli)").assign(Fit_Time=fit_time[5]))
summary_df = summary_df.append(ncentroid_df.assign(Model="Centroid").assign(Fit_Time=fit_time[6]))
summary_df = summary_df.append(mlp_df.assign(Model="Multi Layer Perceptron").assign(Fit_Time=fit_time[7]))

summary_df = summary_df[summary_df['Settings']=='Optimized'].drop('Settings', axis=1).reset_index(drop=True)[["Model", "Data", "Accuracy %", "Precision %", "Recall %", "F1 %", "Prediction Time (s)", "Fit_Time"]].rename({"Fit_Time": "Train Time (s)"}, axis=1)
summary_df
```

Out[95]:

	Model	Data	Accuracy %	Precision %	Recall %	F1 %	Prediction Time (s)	Train Time (s)
0	Decision Tree	Train	72.02	74.65	67.25	70.75	37.84	1.041455
1	Decision Tree	Validate	71.71	74.77	66.61	70.45	14.69	1.041455
2	Decision Tree	Test	71.66	74.18	66.76	70.28	15.62	1.041455
3	Random Forest	Train	86.24	86.22	86.49	86.36	97.95	29.303271
4	Random Forest	Validate	74.58	75.27	74.19	74.72	38.36	29.303271
5	Random Forest	Test	74.60	74.88	74.32	74.60	37.50	29.303271
6	SVC	Train	62.15	68.22	46.45	55.27	54.91	0.780528
7	SVC	Validate	62.14	68.57	46.61	55.49	22.26	0.780528
8	SVC	Test	62.06	67.77	46.55	55.19	25.12	0.780528
9	Gradient Boosting	Train	73.14	73.55	72.83	73.19	91.03	53.543317
10	Gradient Boosting	Validate	72.68	73.43	72.17	72.79	31.50	53.543317
11	Gradient Boosting	Test	72.69	73.06	72.22	72.64	26.56	53.543317
12	Bagging	Train	90.74	92.49	88.82	90.62	103.89	114.754657
13	Bagging	Validate	71.77	73.36	69.49	71.37	37.63	114.754657
14	Bagging	Test	71.89	73.11	69.59	71.31	34.34	114.754657
15	Naive Bayes (Bernoulli)	Train	58.85	58.47	63.00	60.65	43.24	0.478188
16	Naive Bayes (Bernoulli)	Validate	58.76	58.64	63.02	60.75	17.19	0.478188
17	Naive Bayes (Bernoulli)	Test	58.42	57.92	62.71	60.22	17.18	0.478188
18	Centroid	Train	71.81	74.58	66.75	70.45	38.11	0.432970
19	Centroid	Validate	71.51	74.72	66.11	70.15	15.74	0.432970
20	Centroid	Test	71.51	74.17	66.34	70.04	15.62	0.432970
21	Multi Layer Perceptron	Train	71.79	71.28	73.61	72.43	51.69	115.882649
22	Multi Layer Perceptron	Validate	71.11	70.87	72.93	71.89	17.19	115.882649
23	Multi Layer Perceptron	Test	71.32	70.72	73.12	71.90	18.75	115.882649

In []:

In []:

In []: